

Our solution will be similar to the algorithm for VERTEX COVER from the first section of this chapter. Consider the notation defined in the problem. For an element $a \in A$, we *reduce the instance by a* by deleting a from A , and deleting all sets B_i that contain a . Thus, reducing the instance by a producing a new, presumably smaller, instance of HITTING SET.

We observe the following fact. Let $B_i = \{x_1, \dots, x_c\} \subseteq A$ be any of the given sets in the HITTING SET instance. Then at least one of x_1, \dots, x_c must belong to any hitting set H . So by analogy with (2.3) from the notes, we have the following fact

- Let $B_i = \{x_1, \dots, x_c\}$ There is k -element hitting set for the original instance if and only if, for some $i = 1, \dots, c$, the instance reduced by x_i has a $(k - 1)$ -element hitting set.

The proof is completely analogous to that of (2.3). If H is a k -element hitting set, then some $x_i \in H$, and so $H - \{x_i\}$ is a $(k - 1)$ -element hitting set for the instance reduced by x_i . Conversely, if the instance reduced by x_i has a $(k - 1)$ -element hitting set H' , then $H' \cup \{x_i\}$ is a k -element hitting set for the original instance.

Thus, our algorithm is as follows. We pick any set $B_i = \{x_1, \dots, x_c\}$. For each x_i , we recursively test if the instance reduced by x_i has a $(k - 1)$ -element hitting set. We return “yes” if and only if the answer to one of these recursive calls is “yes.” Our running time satisfies $T(m, k) \leq cT(m, k - 1) + O(cm)$, and so it satisfies $T(m, k) = O(c^k \cdot kcm)$. This gives the desired bound, with $f(c, k) = kc^{k+1}$ and $p(m) = m$.

¹ex579.588.787

(a) We prove this by induction on d . If $d = 0$, then Φ is a satisfying assignment, and $Explore(\Phi, d)$ returns “yes.”

Now consider $d > 0$. If $Explore(\Phi, d)$ returns “yes,” it is because one of the recursive calls $Explore(\Phi_i, d - 1)$ returns “yes”; by induction, this means that Φ_i has distance $d - 1$ to a satisfying assignment, and so Φ has distance d to a satisfying assignment.

Conversely, suppose Φ has distance d to a satisfying assignment Φ' . Consider any clause unsatisfied by Φ ; since Φ' satisfies it, it must disagree with Φ on the setting of at least one of the variables in this clause. Thus, one of the assignments Φ_i , which changes the assignment to this variable, is at distance $d - 1$ to Φ' ; by induction the recursive call $Explore(\Phi_i, d - 1)$ will return “yes,” and so the full call $Explore(\Phi, d)$ will also return “yes.”

The running time for $Explore$ satisfies the recurrence $T(n, d) \leq 3T(n, d - 1) + p(n)$, for a polynomial p . Unwinding this to get d down to 0, we have a running time of $O(3^d \cdot p(n))$.

(b) We let Φ_0 denote the assignment in which all variables are set to 0, and we let Φ_1 denote the assignment in which all variables are set to 1. If there is any satisfying assignment, it is within distance at most $n/2$ of one of these, so we can call both $Explore(\Phi_0, n/2)$ and $Explore(\Phi_1, n/2)$, and see if either of these returns “yes.”

The running time of each of these calls is $O(p(n) \cdot 3^{n/2}) = O(p(n) \cdot (\sqrt{3})^n)$.

¹ex695.88.327

Consider an ordered triple (S, i, j) , $1 \leq i, j \leq n$ and S is a subset of the vertices that includes v_i and v_j . Let $B[S, i, j]$ denote the answer to the question, “Is there a Hamiltonian path on $G[S]$ that starts at v_i and ends at v_j ?” Clearly, we are looking for the answer to $B[V, 1, n]$.

We now show how to construct the answers to all $B[S, i, j]$, starting from the smallest sets and working up to larger ones, spending $O(n)$ time on each. Thus the total running time will be $O(2^n \cdot n^3)$.

$B[S, i, j]$ is true if and only if there is some vertex $v_k \in S - \{v_i\}$ so that (v_i, v_k) is an edge, and there is a Hamiltonian path from v_k to v_j in $G[S - \{v_i\}]$. Thus, we set $B[S, i, j]$ to be true if and only if there is some $v_k \in S - \{v_i\}$ for which $(v_i, v_k) \in E$ and $B[S - \{v_i\}, k, j]$ is true. This takes $O(n)$ time to determine.

We claim that such a graph G has a tree decomposition $(T, \{V_t\})$ in which each piece V_t corresponds uniquely to an internal triangular face of G . We prove this by induction on the number of nodes in G .

Choose any internal edge $e = (u, v)$ of G ; deleting u and v produces two components A and B . Let G_1 be the subgraph induced on $A \cup \{u, v\}$ and G_2 the subgraph induced on $B \cup \{u, v\}$. By induction, there are tree decompositions $(T_1, \{X_t\})$ and $(T_2, \{Y_t\})$ of G_1 and G_2 respectively in which the pieces correspond uniquely to internal faces. Thus there are nodes $t_1 \in T_1$ and $t_2 \in T_2$ that correspond to the faces containing the edge (u, v) . If we let T denote the tree obtained by adding an edge (t_1, t_2) to $T_1 \cup T_2$, then $(T, \{X_t\} \cup \{Y_t\})$ is a tree decomposition having the desired properties.

(a) Let us root the tree at an arbitrary node r , and define subtrees T_u as we have done in Chapter 10.2 when solving the Weighted Independent Set Problem. There we defined two subproblems corresponding to each subtree depending on whether or not we include the root u in the set. We will use the same subproblems: $OPT_{in}(u)$ denotes the maximum weight of an independent set of T_u that includes u , and $OPT_{out}(u)$ denotes the maximum weight of an independent set of T_u that does not include u . Now the optimum we are looking for is $\min(OPT_{in}(r), OPT_{out}(r))$. It helps us to define a third subproblem for each subtree: $OPT_{un}(u)$ denotes the maximum weight of an independent set of T_u that does not have to dominate u . When u 's parent is included in the dominating set, u is already dominated by its parent, hence the set selected in the subtree T_u does not have to dominate u .

Now that we have our sub-problems, it is not hard to see how to compute these values recursively. For a leaf $u \neq r$ we have that $OPT_{out}(u) = \infty$, $OPT_{in}(u) = c(u)$, and $OPT_{un}(u) = 0$. For all other nodes u we get a recurrence that defines $OPT_{out}(u)$, $OPT_{in}(u)$, and $OPT_{un}(u)$ using the values for u 's children.

(1) For a node u , the following recurrence defines the values of the sub-problems:

- $OPT_{in}(u) = c(u) + \sum_{v \in \text{children}(u)} OPT_{un}(v)$
- $OPT_{un}(u) = \sum_{v \in \text{children}(u)} \min(OPT_{in}(v), OPT_{out}(v))$
- $OPT_{out}(u) = \min_{v \in \text{children}(u)} (OPT_{in}(v) + \sum_{w \in \text{children}(u), w \neq v} \min(OPT_{out}(v), OPT_{in}(v)))$.

Using this recurrence, we get a dynamic programming algorithm by building up the optimal solutions over larger and larger sub-trees. We define arrays $Mo[u]$, $Mi[u]$ and $Mu[u]$, which hold the values $OPT_{out}(u)$, $OPT_{in}(u)$ and $OPT_{un}(u)$ respectively. For building up solutions, we need to process all the children of a node before we process the node itself.

To find a minimum-weight dominating set of a tree T :

Root the tree at a node r .

For all nodes u of T in post-order

 If u is a leaf then set the values:

$$\begin{aligned} Mo[u] &= \infty \\ Mi[u] &= c(u) \\ Mu[u] &= 0 \end{aligned}$$

 Else set the values:

$$\begin{aligned} Mi[u] &= c(u) + \sum_{v \in \text{children}(u)} Mu[v]. \\ Mu[u] &= \sum_{v \in \text{children}(u)} \min(Mi[v], Mo[v]). \end{aligned}$$

¹ex573.411.14

$$Mo[u] = \min_{v \in \text{children}(u)} Mi[v] + \sum_{w \in \text{children}(u), w \neq v} \min(Mo[v], Mi[v])$$

Endif
Endfor
Return $\max(Mo[r], Mi[r])$.

The algorithms clearly runs in polynomial time, as there are $3n$ subproblems for an n node tree, and each value associated with a subproblem of u of degree n_u can be determined in $O(n_u)$ time. So the total time is $O(\sum_u n_u) = O(n)$.

(b) We extend the algorithm for the case of bounded tree-width by having subproblems associated with the nodes of the tree-decomposition. For each node t of the tree-decomposition, let V_t be the subset of nodes corresponding to tree-node t , T_t the subtree rooted at t , and G_t the corresponding subgraph. Now for each disjoint sets $U, W \subset V_t$ we will define a subproblem and have $OPT(t, U, W)$ the minimum weight of a set in G_t that contains exactly the nodes U in V_t and covers all nodes of G_t except possibly not dominating the a subset of W . Recall that in the case of a tree, the recurrence for $OPT_{out}(u)$ was a little awkward, as we needed to select a child of u that is covering u . Here we would need to do this for each node in $V_t - (U \cup W)$. To make this simpler to write, we will define more subproblems. Let t be a node of the tree decomposition, and let t_1, \dots, t_d be its children, then we define subproblems $OPT(t, i, U, W)$ for each $0 \leq i \leq d$ to be the minimum weight of a set in the graph corresponding to the union of subtrees T_{t_1}, \dots, T_{t_i} and the set V_t that contains exactly the nodes U in V_t and dominates all nodes of this subgraph except possibly not dominating the a subset of W .

So if d_r is the degree of the root, then the optimum value we are looking for is $\min_{U \subset V_r} OPT(r, d_r, U, \emptyset)$. For any node t we have $OPT(t, 0, U, W) = \sum_{u \in U} c(u)$ if U dominates the nodes $V_t - W$, and ∞ otherwise. This defines the values at the leafs.

Given the subproblems, we will get the value at a node t using the values for smaller i , and the values at the children of t as follows. For a set U we will use the notation $c(U) = \sum_{u \in U} c(u)$, and $\delta(U)$ is the set dominated by U . Let t be a node of the tree decomposition and t_1, \dots, t_d its children, let n_i be the degree of child i .

(2) The value of $OPT(t, i, U, W)$ for $i \geq 1$ is given by the following recurrence:

$$OPT(t, i, U, W) = c(U) + \min_{U_i \subseteq V_{t_i}: U_i \cap V_t = U \cap V_{t_i}} (OPT(t, i-1, U, W \cup \delta(U_i)) + OPT(t_i, n_i, U_i, (\delta(U) \cup W) \cap V_{t_i}) - c(U \cap V_{t_i}))$$

For a tree-decomposition of width k there are 3^{k+1} subproblems associated with a (t, i) pair, and there are n such pairs if the tree is of size n . Computing each value takes only $O(1)$ time, so the total time required is $O(3^k n)$.

Let $(T; \{V_t | t \in T\})$ be the given tree decomposition rooted at r . There are k $(s_i; t_i)$ terminal pairs. We focus on the tree-width 2 case. For convenience, we assume that there are no two pieces V_{t_1} and V_{t_2} where (t_1, t_2) is an edge and $V_{t_1} \subset V_{t_2}$. Consider the subgraph G_t . Note that there can be at most one i for which P_i both enters and leaves G_t , since any such path uses up at least 2 vertices of V_t . Note also that there can be at most 3 s_i-t_i terminal pairs that have one end in G_t and the other one outside (as the paths connecting such pairs must go through V_t).

- If there are 3 such pairs, that each node $v \in V_t$ must be connected via disjoint paths to one of them, and terminal pairs inside G_t must connect via paths inside G_t . There are $O(1)$ cases here to consider depending on which of the nodes in V_t is used to connect which of the 3 separated terminal pairs.
- If there are 2 such terminal pairs, than of the at most 3 nodes in V_t 2 must be connected via disjoint paths to one of them, the third is either not used in any of the paths or is used by path connecting two terminals s_i and t_i inside, or two terminals outside G_t . Now there are $O(k)$ cases to consider, depending on which terminal pair i is using the extra node.
- If there is only one such pair, than we can have one pair i with both terminals inside or outside of G_t , that uses 2 nodes in V_t , or the one path leaving G_t , can leave, come back and leave again, or one or two paths can use just one node in V_t while having both terminals inside or both outside of G_t . There are $O(k^2)$ cases to consider here.
- If there are no such pairs, than one path can use 2 or 3 nodes in V_t , or multiple paths can use one node each. Now there are $O(k^3)$ cases to consider.

We define multiple subproblems for each t according to the possibilities discussed above. For the at most 3 nodes in V_t there are at most $O(k^3)$ possible cases. This defines $O(k^3n)$ subproblems. The value of a subproblem is simply 0 or 1 (or true or false) depending whether or not there are disjoint paths in G_t that satisfy the state of the nodes in V_t corresponding to the subproblem, that is, connect each $v \in V_t$ to the terminal in question inside G_t (and possibly connect the two nodes in V_t to each other, if needed), via disjoint paths inside G_t . The desired disjoint paths exists if and only if the value of one of the subproblems that connects all terminal pairs within the subgraph G_r (which is the whole graph).

Given values for all the subproblems associated with the children t_1, \dots, t_d of a node t , we want to get the value of the given subproblem efficiently. To do this consider a node $v \in V_t$, the subproblem under question wants a particular paths P_i to go through this vertex in $O(d)$ time.

¹ex209.650.476

Checking whether G is 1- or 2-colorable is easy. For $k = 3, 4, \dots, w + 1$, we test whether G is k -colorable by dynamic programming. We use notation similar to what we used in the Maximum-Weight Independent Set problem for graphs of bounded tree-width. Let $(T, \{V_t : t \in T\})$ be a tree decomposition of G . For the subtree rooted at t , and every coloring χ of V_t using the color set $\{1, 2, \dots, k\}$, we have a predicate $q_t(\chi)$ that says whether there is a k -coloring of G_t that is equal to χ when restricted to V_t . This requires us to maintain $k^{(w+1)} \leq (w+1)^{(w+1)}$ values for each piece of the tree decomposition.

We compute the values $q_t(\chi)$ when t is a leaf by simply trying all possible colorings of G_t . In general, suppose t has children t_1, \dots, t_d , and we know the values of $q_{t_i}(\chi)$ for each choice of t_i and χ . Then there is a coloring of G_t consistent with χ on V_t if and only if there are colorings of the subgraphs G_{t_1}, \dots, G_{t_d} that are consistent with χ on the parts of V_{t_i} that intersect with V_t . Thus we set $q_t(\chi)$ equal to *true* if and only if there are colorings χ_i of V_{t_i} such that $q_{t_i}(\chi_i) = \text{true}$ and χ_i is the same as χ when restricted to $V_t \cap V_{t_i}$.

¹ex897.854.812

We build the following bipartite graph G . There is a node u_i for each variable x_i , a node v_j for each clause C_j , and an edge (u_i, v_j) whenever x_i appears in C_j .

We first note that since each variable appears three times, and each clause has length three, the number of nodes on the left side of G equals the number of nodes on the right side. More strongly, we claim that G in fact has a perfect matching.

This is a consequence of a more general claim: that if every node in a bipartite graph G has the same degree d , then G has a perfect matching. (Here $d = 3$.) Indeed, if G does not have a perfect matching, then by Hall's Theorem it has a set A on the left side for which $|A| < |\Gamma(A)|$, where $\Gamma(A)$ denotes the set of neighbors of any node in A . But A has $d|A|$ nodes coming out of it, and $\Gamma(A)$ can only absorb $d|\Gamma(A)|$ of them, so this is not possible.

Consider the perfect matching in the graph G we constructed from the 3 -SAT instance. For each variable, we set it in a way that satisfies the clause it is matched to. In this way, we satisfy the full collection of clauses.

¹ex592.206.332

We root the tree at some node r and associate a subtree T_v with each node $v \in T$, and let n_v denote the number of nodes in the subtree T_v . A solution for the graph partitioning problem may split the subtree T_v in any ratio, however, it is not hard to see that if (A, B) is the maximum weight cut splitting the tree T into two equal sides, and assume $v \in A$ and $|A \cap T_v| = k$, then the induced partition of T_v by having $A_v = A \cap T_v$, and $B_v = B \cap T_v$ is the maximum weight cut separating T_v into two sides, where the side containing v has size k . This is true as the only interaction of the partition of T_v with the remaining graph is through node v . More precisely, assume it is false, and consider the optimal such cut (A', B') . Replacing A_v by A' keeps all edges across the cut, except replacing edges that cross the (A_v, B_v) cut with edges that cross the (A', B') cut, and hence it results in a cut of larger weight. This contradiction proves the claim.

Given the above observation, we define subproblems for each subtree T_v and each integer $k \leq n_v$, where $OPT(v, k)$ is the maximum cut of the subgraph T_v separating T_v into two sides where the side containing v has size k . The final answer is $OPT(r, n/2)$ if $n = n_r$ is the number of nodes in T .

We will use $c(u, v)$ as the cost or weight of the edge $e = (v, w)$ of the tree. For a leaf v we have $k = 1$ as the only possible value and $OPT(v, 1) = 0$. Now consider $OPT(v, k)$ for some non-leaf v . The side A containing v must contain $k - 1$ nodes in addition to node v . If v has only one child u then we need to consider two cases depending on whether or not the child u is on the same side of the cut as v , so we get that in this case

$$OPT(v, k) = \max(opt(u, k - 1), c(v, u) + OPT(u, n_u - k + 1)).$$

If v has two children u and w , then we will consider all possible ways of dividing these $k - 1 = \ell_1 + \ell_2$ nodes between the two subtrees rooted at the two children of v . For each such division, there are further cases depending whether or not the root of these subtrees is on the same side of the cut as v . We get the following recurrence in this case

$$\begin{aligned} OPT(v, k) = & \max_{0 \leq \ell_1 \leq k-1, \ell_2 = k-1-\ell_1} (\max(OPT(u, \ell_1), c(v, u) + OPT(u, n_u - \ell_1)) \\ & + \max(OPT(w, \ell_2), c(v, w) + OPT(w, n_w - \ell_2))). \end{aligned}$$

There are $O(n^2)$ subproblems. We can compute the values of the subproblems starting at the leaves, and gradually considering bigger subtrees. This recurrence allows us to compute the value of a subproblem in $O(1)$ time given the values of the subproblems associated with smaller trees. So the total time required is $O(n^2)$.