



## COMP3511 Operating Systems

### Topic 4: Multithreaded Programming

Dr. Desmond Tsoi

Department of Computer Science & Engineering  
The Hong Kong University of Science and Technology  
Hong Kong SAR, China



Acknowledgment: The lecture notes are based on various sources on the Internet

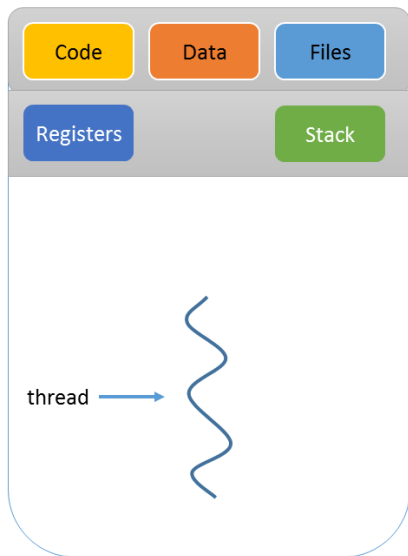
# Threads

## Threads

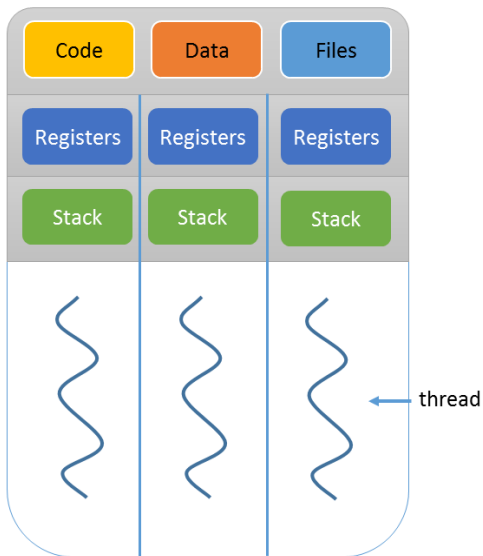
A **thread** is short for a **thread of execution**. A thread is a **portion of code that may be executed independently of the main program**.

- Most modern applications and/or programs are multithreaded
- Threads **run within an application or a process**
- Multiple tasks with the application can be implemented by separate threads, e.g., a word processor may have separate threads:
  - ▶ Render graphics
  - ▶ Respond to keystrokes from user
  - ▶ Spell/grammar checking
- Process creation is heavy-weight while **thread creation is light-weight**
- Can simplify code, increase efficiency
- **Kernels** are generally **multithreaded**

# Single and Multithreaded Processes

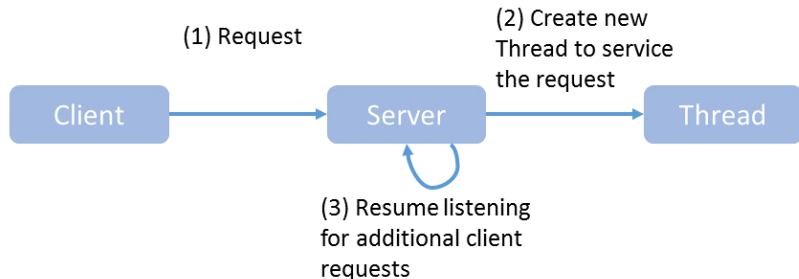


Single-threaded process



Multithreaded process

# Multithreaded Server Architecture



- A **single application** may be required to **perform several similar tasks**. For example a busy web server may process thousands of web requests concurrently. Creating one process for each client request is cumbersome (resource-intensive) and time-consuming
- A **single application** may need to **do multiple tasks**. For example, a web browser (client) needs to display images or text (one thread) while another thread retrieves data from the network

# Benefits

- Responsiveness

- ▶ May allow continued execution if part of process is blocked, especially important for user interface

- Resource sharing

- ▶ Threads within a process share resources of the process by default, easier than shared memory or message passing that must be explicitly arranged by the programmer

- Economy


- ▶ Thread creation is much cheaper than process creation, thread switching also has much lower overhead than context switching (switching to a different process)

- Scalability

- ▶ A process can take advantage of multiprocessor architectures by running multiple threads of the process simultaneously on different processors (CPUs)

# Multicore Programming

- Multicore or multiprocessor systems put pressure on programmers to make better use of the multiple computing cores
- **Programming challenges in multicore systems** include:
  - ▶ **Identifying tasks**: To divide applications into separate, concurrent tasks
  - ▶ **Balance**: Tasks perform equal work of equal value
  - ▶ **Data splitting**
  - ▶ **Data dependency**
  - ▶ **Testing and debugging**



# Parallelism vs. Concurrency


- **Parallelism** implies a system can perform more than one task simultaneously
- **Concurrency** supports more than one task making progress
  - ▶ Single processor / core, scheduler providing concurrency
- **Types of parallelism**
  - ▶ **Data parallelism** - distributes subsets of the same data across multiple cores, same operation on each
  - ▶ **Task parallelism** - distributing threads across cores, each thread performing unique operation




Poppy Cat

# Parallelism vs. Concurrency (Cont'd)

- Concurrent execution on single-core system:



- Parallelism on a multicore system





# Amdahl's Law

- Identifies performance gains from adding additional cores to an application that has both serial and parallel components
- S denotes percentage of serial portion and N denotes number of processing cores

$$speedup \leq \frac{1}{S + \frac{1-S}{N}}$$

- Example: What is the speedup if an application is 75% parallel and 25% serial, moving from 1 to 2 cores?

$$speedup \leq \frac{1}{0.25 + \frac{1-0.25}{2}} = 1.6$$

- What if N approaches infinity?

$$speedup \leq \frac{1}{S + \frac{1-S}{\infty}} = 1/S$$


- What if N is 1, i.e. single core?

- Serial portion of an application has disproportionate effect on performance gained by adding additional cores

# Thread State

- Each thread has a **Thread Control Block (TCB)**
  - ▶ **Execution information**: CPU registers, program counter, pointer to stack
  - ▶ **Scheduling information**: State (more later), priority, CPU time
  - ▶ **Accounting information**
  - ▶ **Various pointers** (for implementing scheduling queues)
- OS keeps track of TCBs in protected memory
  - ▶ Array / linked list, ...
- **Information shared by all threads in process / address space**
  - ▶ **Contents of memory** (global variables, heap)
  - ▶ **I/O state** (file system, network connections, etc.)
- **State "private" to each thread**
  - ▶ Kept in TCB
  - ▶ CPU registers (including program counter)
  - ▶ Execution stack
    - ★ **Parameters, temporary variables**
    - ★ Keep **program counters** while called procedures are executing

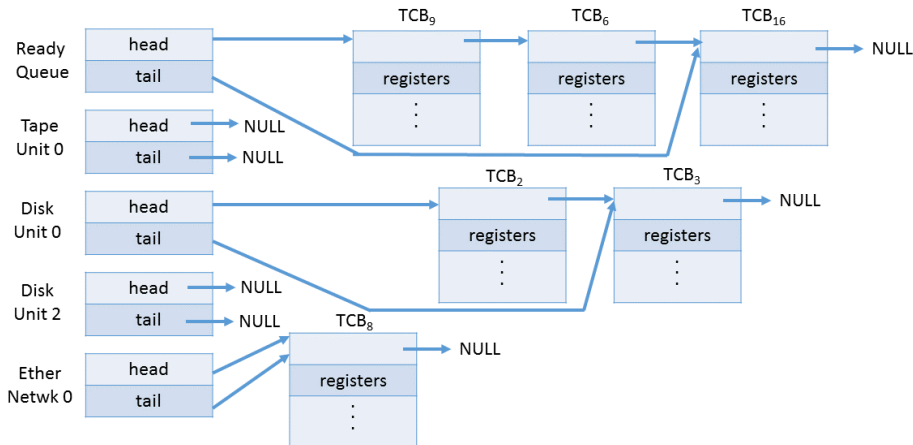
# Life Cycle of a Thread



- As a thread executes, it changes **state**
  - ▶ **new**: The thread is being created
  - ▶ **ready**: The thread is waiting to run
  - ▶ **running**: Instructions are being executed
  - ▶ **waiting**: Thread waiting for some event to occur
  - ▶ **terminated**: The thread has finished execution
- "Active" threads are represented by their TCBs
  - ▶ TCBs organized into queues based on their states

# Ready Queue and Various I/O Device Queues

- Thread not running → TCB is in some scheduler queue
  - ▶ Separate queue for each device/signal/condition
  - ▶ Each queue can have a different scheduler policy



# Examples of Multithreaded Programs


- Embedded systems
  - ▶ Elevators, Planes, Medical Systems, Wristwatches
- Most modern OS kernels
  - ▶ Internally concurrent to deal with concurrent requests by multiple users
  - ▶ But no protection needed within kernel
- Database servers
  - ▶ Access to shared data by many concurrent users
  - ▶ Also background utility programming must be done
- Network servers
  - ▶ Concurrent requests from network
  - ▶ Again, single programs, multiple concurrent operations
  - ▶ File server, Web server, the airline reservation systems
- Parallel programming (More than one physical CPU)
  - ▶ Split program into multiple threads for parallelism

# User Threads and Kernel Threads

- **Support for threads** may be provided at either the **user level, for user threads**, or by the **kernel, for kernel threads**
- **User threads** - **management done by user-level threads library without kernel support**. Three primary thread libraries
  - ▶ POSIX Pthreads
  - ▶ Win32 threads
  - ▶ Java threads
- **Kernel threads** - **supported by the kernel**. Virtually all general-purpose operating systems support kernel threads, including
  - ▶ Windows
  - ▶ Solaris
  - ▶ Linux
  - ▶ Mac OS X
- Ultimately, a **relationship must exist between user threads and kernel threads**
  - ▶ Many-to-One
  - ▶ One-to-One
  - ▶ Many-to-Many


# Many-to-One

- Many user-level threads mapped to a single kernel thread
- One thread blocking causes all to block
- Multiple threads may not run in parallel on a multicore system because only one may be in kernel at a time
- Few systems currently use this model. For examples:
  - ▶ Solaris Green Threads
  - ▶ GNU Portable Threads



# One-to-One


- Each user-level thread maps to kernel thread
- Creating a user-level thread creates a kernel thread
- More concurrency than many-to-one
- Number of threads per process sometimes restricted due to overhead
- Examples:
  - ▶ Windows NT/XP/2000
  - ▶ Linux
  - ▶ Solaris 9 and later






# Many-to-Many

- Allows many user-level threads to be mapped to many kernel threads
- Allows the operating system to create a sufficient number of kernel threads
- Examples
  - ▶ Solaris prior to version 9
  - ▶ Windows NT/2000 with the ThreadFiber package




# Two-Level Model

- Similar to many-to-many, except that it also allows a user thread to be bound to kernel thread
- Examples
  - ▶ IRIX
  - ▶ HP-UX
  - ▶ Tru64 UNIX
  - ▶ Solaris 8 and earlier



# Thread Libraries

- Thread library provides programmer with API for creating and managing threads
- Implementing in two primary ways
  - ▶ Library entirely in user space with no kernel support. This means that invoking a function in the library results in a local function call in user space, and not a system call
  - ▶ Kernel-level library supported directly by the OS
- Three main thread libraries are in use today:
  - ▶ POSIX Pthreads
  - ▶ Windows
  - ▶ Java



# POSIX Pthreads

## ● Create a thread

```
int pthread_create(pthread_t* thread,  
                  const pthread_attr_t* attr,  
                  void* (*start_routine)(void*),  
                  void* arg)
```

- ▶ thread - returns the thread id
- ▶ attr - set to NULL if default thread attributes are used
- ▶ start - pointer to function to be threaded
- ▶ arg - pointer to argument of function
- ▶ return 0 if the thread creation is successful

## ● Waits for the thread to terminate

```
int pthread_join(pthread_t th, void** thread_return)
```

- ▶ th - thread suspended until the thread identified by th terminates
- ▶ thread\_return - if thread\_return is not NULL, the return value of th is stored in the location pointed to by thread\_return
- ▶ return 0 if the thread join is successful

## ● Terminates a thread

```
void pthread_exit(void* retval)
```

- ▶ retval - return value of pthread\_exit()

# POSIX Pthreads: Example

```
/* threadexample.c */
#include <stdio.h>
#include <stdlib.h>
#include <pthread.h>

void* print(void *ptr) { printf("%s \n", (char*)ptr); }

int main() {
    pthread_t t1, t2;
    const char* msg1 = "Thread 1";
    const char* msg2 = "Thread 2";
    int ret1, ret2;

    ret1 = pthread_create(&t1, NULL, print, (void*)msg1);           // create a thread
    if(ret1) { fprintf(stderr, "Error, pthread_create(): %d\n", ret1); exit(1); }
    ret2 = pthread_create(&t2, NULL, print, (void*)msg2);           // create a thread
    if(ret2) { fprintf(stderr, "Error, pthread_create(): %d\n", ret2); exit(1); }

    printf("pthread_create() for thread 1: %d\n", ret1);
    printf("pthread_create() for thread 2: %d\n", ret2);
    pthread_join(t1, NULL);
    pthread_join(t2, NULL);
    exit(0);
}
```

# POSIX Pthreads: Example (Cont'd)

- Command to compile the program:  
`gcc -pthread threadexample.c`
- Output:  
`pthread_create() for thread 1: 0`  
`pthread_create() for thread 2: 0`  
Thread 1  
Thread 2



# Threading Issues - Semantics of `fork()` and `exec()` in Multithreaded Environment

- Does `fork()` duplicate only the calling thread or all threads?
  - ▶ Some UNIX have two versions of `fork`, one version of `fork()` duplicates all threads and the other duplicates only the thread that invoked the `fork()` system call
  - ▶ If `exec()` is called immediately after forking, then duplicating all threads is unnecessary, as the program specified in the parameters to `exec()` will replace the process
- `exec()` usually works as normal
  - ▶ Replace the running process including all threads

# Signal in UNIX

## Signals

Signals are in UNIX systems to notify a process that a particular event has occurred

- A signal may be received either synchronously or asynchronously, depending on the source of and the reason for the event being signaled
- All signals follow the same pattern
  - ▶ Signal is generated by the occurrence of a particular event
  - ▶ Signal is delivered to a process
  - ▶ Once delivered, the signal must be handled
- A signal may be handled by one of the two possible handlers
  - ▶ A default signal handler
  - ▶ A user-defined signal handler
- Every signal has default handler that kernel runs when handling signal
  - ▶ User-defined signal handler can override the default handler
  - ▶ For single-threaded, signal delivered to process



# UNIX / Linux Signal Handling

- Prototype of a signal handling function

```
void <signal handler func name>(int sig)
```

- To get the **signal handler function registered to the kernel**, the signal handler function pointer is passed as **second argument** to the 'signal' function. The prototype of the signal function is

```
void (*signal(int signo, void (*func )(int)))(int);
```

where signo refers to signal number

- The following shows some of the **signal numbers**
  - ▶ SIGHUP: Hangup, report that user's terminal is disconnected
  - ▶ SIGTERM: Termination, it can be blocked, handled, and ignored. Generated by "kill" command
  - ▶ SIGINT: Interrupt, Program interrupt signal from keyboard (normally Ctrl-c)
  - ▶ SIGQUIT: Quit, terminate process and generate core dump
  - ▶ SIGFPE: Floating point arithmetic exception, e.g. division by zero
  - ▶ SIGKILL: Kill, unblockable, cause immediate program termination, cannot be handled, blocked or ignored
  - ▶ SIGCHLD: Child status has changed, child process stopped or terminated
  - ▶ SIGSEGV: Segmentation violation, dereferencing a bad or NULL pointer

# UNIX / Linux Signal Handling: Example 1

```
#include <unistd.h>
#include <stdio.h>
#include <stdlib.h>
#include <signal.h>

// Define the function to be called when ctrl-
c
// (SIGINT) signal is sent to process
void signal_callback_handler(int signum) {
    printf("Caught signal %d\n",signum);
    // Terminate program
    exit(1);
}

int main() {
    // register signal and signal handler
    signal(SIGINT, signal_callback_handler);

    while(1) {
        printf("Processing here.\n");
        sleep(1);
    }
    return 0;
}
```

- Run the program and press Ctrl-c:  
Processing here.  
Processing here.  
Processing here.  
Caught signal 2
- Run the program at one terminal and type a command at another terminal
  - ▶ **First terminal running the program**  
Processing here.  
Processing here.  
Processing here.  
Caught signal 2
  - ▶ **Second terminal (Type command)**  
**Assume the process id of the program is 12570 (check using ps -all)**  
kill -INT 12570

## UNIX / Linux Signal Handling: Example 2

*// Example of how two processes can  
communicate to each other using  
kill() and signal(). We will fork()  
2 process and let the parent  
send a few signals to it's child*

```
#include <stdio.h>
#include <stdlib.h>
#include <signal.h>
```

```
void sighup();
void sigint();
void sigquit();
```

```
int main() {
    int pid;

    if((pid = fork()) < 0) {
        perror("Fail to create child");
        exit(1);
    }
}
```

```
if(pid == 0) {                                // child
    // register functions to the kernel
    signal(SIGHUP,sighup);
    signal(SIGINT,sigint);
    signal(SIGQUIT, sigquit);
    for(;;);                                  // infinite loop
}
else {                                          // parent
    printf("\nParent:  send SIGHUP\n\n");
    kill(pid,SIGHUP);
    sleep(3);                                // pause for 3 seconds
    printf("\nParent:  send SIGINT\n\n");
    kill(pid,SIGINT);
    sleep(3);                                // pause for 3 seconds
    printf("\nParent:  send SIGQUIT\n\n");
    kill(pid,SIGQUIT);
    sleep(3);                                // pause for 3 seconds
}
```

## UNIX / Linux Signal Handling: Example 2 (Cont'd)

```
void sighup() {  
    signal(SIGHUP,sighup);           // reset signal  
    printf("Child: I have received a SIGHUP\n");  
}  
  
void sigint() {  
    signal(SIGINT,sigint);           // reset signal  
    printf("Child: I have received a SIGINT\n");  
}  
  
void sigquit() {  
    printf("My parent has killed me!!!\n");  
    exit(0);  
}
```

### Output:

Parent: send SIGHUP

Child: I have received a SIGHUP

Parent: send SIGINT

Child: I have received a SIGINT

Parent: send SIGQUIT

My parent has killed me!!!

# Threading Issues - Signal in UNIX (Cont'd)

- Where should a signal be delivered a multi-threaded program?
  - ▶ Deliver the signal to the thread to which the signal applies
  - ▶ Deliver the signal to every thread in the process
  - ▶ Deliver the signal to certain threads in the process
  - ▶ Assign a specific thread to receive all signals for the process
- The method for delivering a signal depends on the type of signal
  - ▶ Synchronous signals need to be delivered to the thread causing the signal, not other threads
  - ▶ Terminating a process signal should be sent to all threads with the process

© MARK ANDERSON

WWW.ANDERSTOONS.COM



"Normally I can see both sides of an issue. But this..."

# Threading Issues - Thread Cancellation

- Thread cancellation involves terminating a thread before it has completed.

For example: Multiple threads are concurrently searching through a database, one thread returns the result, the remaining threads might be canceled

- Thread to be canceled is called target thread
- Cancellation of a target thread may occur in two different scenarios
  - ▶ Asynchronous cancellation terminates the target thread immediately
  - ▶ Deferred cancellation allows the target thread to periodically check if it should be canceled

© MAZIK ANDERSON

WWW.ANDERSTOONS.COM



"Normally I can see both sides of an issue. But this..."

# Thread-Local Storage

- Thread-local storage (TLS) allows **each thread to have its own copy of data**
- Useful when you do not have control over the thread creation process
- Different from local variables
  - ▶ Local variables visible only during single function invocation
  - ▶ **TLS visible across function invocations**, i.e. similar to static data
- TLS is **unique to each thread**



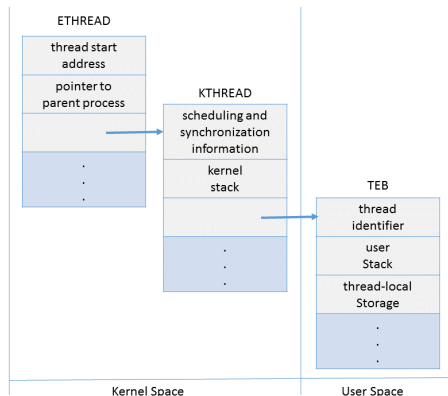
# Windows Threads

- Windows implements the Windows API for Windows 98, NT, 2000, WinXP, and Windows 7
- Implements the one-to-one mapping, one user-level thread → one kernel thread
- Each thread contains
  - ▶ A thread ID uniquely identifying the thread
  - ▶ A register set representing the status of the processor
  - ▶ Separate user and kernel stacks for when thread runs in user mode or kernel mode
  - ▶ A private data storage area used by run-time libraries and dynamic link libraries (DLLs)
- The register set, stacks, and private storage area are known as the context of the thread



# Windows Threads (Cont'd)

- The primary data structures of a thread include
  - ▶ ETHREAD (executive thread block), which includes pointer to process to which thread belongs and to KTHREAD, in kernel space
  - ▶ TEB (thread environment block), which has thread ID, user-mode stack, thread-local storage, in user space
  - ▶ KTHREAD (kernel thread block), which has scheduling and synchronization information, kernel-mode stack, pointer to TEB, in kernel space



# Linux Threads

- Linux refers to processes and threads as tasks rather than threads
- Thread creation is done through `clone()` system call
- `clone()` allows a child tasks to determine how to share the address space of the parent task (process)
  - ▶ Flags control behavior

flag	meaning
CLONE_FS	File-system information is shared
CLONE_VM	The same memory space is shared
CLONE_SIGNAND	Signal handlers are shared
CLONE_FILES	The set of open files is shared

That's all!

Any questions?

